

CS 433 (2) HW5 Report

Boris Nikulin

2019-05-11

Contents

1 Data Analysis	1
1.1 Data Import	1
1.2 Data Analysis	2
1.3 Data Visualization	6

1 Data Analysis

1.1 Data Import

First we load the simulation data, filter out processes that never finished, and give the data a once over.

However, using `readr::read_tsv` to load the data take's minutes because the data per run is around 1 GB. After all the runs of the simulation, the final dataset is 4.2 GB. As a side note, the gzip comprised size of the whole data set is 444 MB. For this reason, we will primarily use `data.table` instead of `readr` and `dplyr`. `data.table::fread` reads the whole dataset in 10 s to 20 s.

```
library(data.table)
library(dplyr)
library(magrittr)

data <- fread('paging_sim.tsv')

data %<>% .[, time := time / 1E6]

glimpse(data)
## Observations: 119,364,413
## Variables: 7
## $ time          <dbl> 0.0e+00, 1.0e-06, 1.1e-05, 1.2...
## $ num_pages     <int> 524288, 524288, 524288, 524288...
## $ page_size_bytes <int> 256, 256, 256, 256, 256, ...
```

```
## $ memory_size_bytes      <int> 65536, 65536, 65536, 65536, 65...
## $ replacement_algorithm <chr> "fifo", "fifo", "fifo", "fifo"...
## $ event                  <chr> "reference", "fault", "referen...
## $ page                   <int> 397863, 397863, 397863, 281415...
```

The times are in seconds.

1.2 Data Analysis

```
run_group <- c(
  'num_pages',
  'page_size_bytes',
  'memory_size_bytes',
  'replacement_algorithm'
)

(data_time <- data[, .(time_max = max(time)), mget(run_group))])

##      num_pages page_size_bytes memory_size_bytes
## 1:      524288           256           65536
## 2:      524288           256           65536
## 3:      524288           256           65536
## 4:      524288           256           65536
## 5:      524288           256          16777216
## 6:      524288           256          16777216
## 7:      524288           256          16777216
## 8:      524288           256          16777216
## 9:       16384          8192           65536
## 10:      16384          8192           65536
## 11:      16384          8192           65536
## 12:      16384          8192           65536
## 13:      16384          8192          16777216
## 14:      16384          8192          16777216
## 15:      16384          8192          16777216
## 16:      16384          8192          16777216
##      replacement_algorithm  time_max
## 1:                fifo 13.591457
## 2:          rng:mt19937  6.234164
## 3:      rng:minstd_rand  6.280066
## 4:                lru 15.925843
## 5:                fifo  6.760374
## 6:          rng:mt19937  6.303923
## 7:      rng:minstd_rand  6.264671
## 8:                lru  9.071865
## 9:                fifo 14.191701
## 10:          rng:mt19937  5.529430
```

```

## 11:      rng:minstd_rand  5.545736
## 12:      lru 16.382787
## 13:      fifo  6.149082
## 14:      rng:mt19937  5.564538
## 15:      rng:minstd_rand  5.540253
## 16:      lru  7.910705

(data_event_count <- data[, .N, mget(append(run_group, 'event'))])

##      num_pages page_size_bytes memory_size_bytes
##  1:    524288          256          65536
##  2:    524288          256          65536
##  3:    524288          256          65536
##  4:    524288          256          65536
##  5:    524288          256          65536
##  6:    524288          256          65536
##  7:    524288          256          65536
##  8:    524288          256          65536
##  9:    524288          256          65536
## 10:    524288          256          65536
## 11:    524288          256          65536
## 12:    524288          256          65536
## 13:    524288          256          65536
## 14:    524288          256          65536
## 15:    524288          256          65536
## 16:    524288          256          65536
## 17:    524288          256         16777216
## 18:    524288          256         16777216
## 19:    524288          256         16777216
## 20:    524288          256         16777216
## 21:    524288          256         16777216
## 22:    524288          256         16777216
## 23:    524288          256         16777216
## 24:    524288          256         16777216
## 25:    524288          256         16777216
## 26:    524288          256         16777216
## 27:    524288          256         16777216
## 28:    524288          256         16777216
## 29:    524288          256         16777216
## 30:    524288          256         16777216
## 31:    524288          256         16777216
## 32:    524288          256         16777216
## 33:    16384          8192          65536
## 34:    16384          8192          65536
## 35:    16384          8192          65536
## 36:    16384          8192          65536

```

## 37:	16384	8192	65536
## 38:	16384	8192	65536
## 39:	16384	8192	65536
## 40:	16384	8192	65536
## 41:	16384	8192	65536
## 42:	16384	8192	65536
## 43:	16384	8192	65536
## 44:	16384	8192	65536
## 45:	16384	8192	65536
## 46:	16384	8192	65536
## 47:	16384	8192	65536
## 48:	16384	8192	65536
## 49:	16384	8192	16777216
## 50:	16384	8192	16777216
## 51:	16384	8192	16777216
## 52:	16384	8192	16777216
## 53:	16384	8192	16777216
## 54:	16384	8192	16777216
## 55:	16384	8192	16777216
## 56:	16384	8192	16777216
## 57:	16384	8192	16777216
## 58:	16384	8192	16777216
## 59:	16384	8192	16777216
## 60:	16384	8192	16777216
## 61:	16384	8192	16777216
## 62:	16384	8192	16777216
## 63:	16384	8192	16777216
## 64:	16384	8192	16777216
##	num_pages	page_size_bytes	memory_size_bytes
##	replacement_algorithm	event	N
## 1:	fifo	reference	5000000
## 2:	fifo	fault	2778912
## 3:	fifo	replace	2778656
## 4:	fifo	flush	2483210
## 5:	rng:mt19937	reference	5000000
## 6:	rng:mt19937	fault	336407
## 7:	rng:mt19937	replace	336151
## 8:	rng:mt19937	flush	61818
## 9:	rng:minstd_rand	reference	5000000
## 10:	rng:minstd_rand	fault	336087
## 11:	rng:minstd_rand	replace	335831
## 12:	rng:minstd_rand	flush	61562
## 13:	lru	reference	5000000
## 14:	lru	fault	2632039
## 15:	lru	replace	2631783

```

## 16:          lru      flush 2336294
## 17:          fifo reference 5000000
## 18:          fifo      fault  514905
## 19:          fifo     replace  449369
## 20:          fifo      flush  300719
## 21:      rng:mt19937 reference 5000000
## 22:      rng:mt19937      fault  332699
## 23:      rng:mt19937     replace  267163
## 24:      rng:mt19937      flush   58969
## 25:  rng:minstd_rand reference 5000000
## 26:  rng:minstd_rand      fault  332344
## 27:  rng:minstd_rand     replace  266808
## 28:  rng:minstd_rand      flush   58657
## 29:          lru reference 5000000
## 30:          lru      fault  482163
## 31:          lru     replace  416627
## 32:          lru      flush  266250
## 33:          fifo reference 5000000
## 34:          fifo      fault  2965735
## 35:          fifo     replace  2965727
## 36:          fifo      flush  2943957
## 37:      rng:mt19937 reference 5000000
## 38:      rng:mt19937      fault   31530
## 39:      rng:mt19937     replace   31522
## 40:      rng:mt19937      flush   16567
## 41:  rng:minstd_rand reference 5000000
## 42:  rng:minstd_rand      fault   31487
## 43:  rng:minstd_rand     replace   31479
## 44:  rng:minstd_rand      flush   16468
## 45:          lru reference 5000000
## 46:          lru      fault  2663297
## 47:          lru     replace  2663289
## 48:          lru      flush  2641522
## 49:          fifo reference 5000000
## 50:          fifo      fault  249038
## 51:          fifo     replace  246990
## 52:          fifo      flush  233736
## 53:      rng:mt19937 reference 5000000
## 54:      rng:mt19937      fault   31497
## 55:      rng:mt19937     replace   29449
## 56:      rng:mt19937      flush   16505
## 57:  rng:minstd_rand reference 5000000
## 58:  rng:minstd_rand      fault   31249
## 59:  rng:minstd_rand     replace   29201
## 60:  rng:minstd_rand      flush   16314

```


```
## 61:          lru reference 5000000
## 62:          lru   fault  213375
## 63:          lru  replace  211327
## 64:          lru   flush  197729
## replacement_algorithm    event      N
data_event_count %<>% .[event != 'reference']
```

Another side note, `awk` takes around 10 s per simulation to generate the counts while `data.table` gets all the counts at once in under 10 s.

1.3 Data Visualization

```
library(ggplot2)

ggplot(data_event_count,
       aes(event, N, fill = replacement_algorithm)) +
  geom_col(position = 'dodge') +
  facet_grid(memory_size_bytes ~ page_size_bytes) +
  labs(x = 'Event', y = 'Count', fill = 'Algorithm')
```



rng:mt19937 refers to the C++ `random` library and is a 32 bit mersenne twister default constructed. rng:minstd_rand also corresponds to the `random` library and is the newer 1993 minimum standard LCG. It is also default constructed.


The bar graph above has some key takeaways. Page size, shown in the horizontal facets, does not reduce the event counts nearly as much as memory size, shown in the vertical facets. At the end of the day, more memory is best.

Another takeaway is that when there is insufficient memory, page size has no effect or a small adverse affect on the number of events excluding the random algorithm. Page size does have an effect when memory size is not the limiting factor.

Lastly, random page replacement is unfairly good for how simple the algorithm is and is primarily affected by page size unlike the other algorithms.

```
ggplot(data_time, aes(replacement_algorithm, time_max)) +
  geom_col() +
```

```
facet_grid(memory_size_bytes ~ page_size_bytes) +
labs(x = 'Algorithm', y = 'Time (s)') +
coord_flip()
```



The times for the final events, which included the time to log each event, are pretty similar.

I expected LRU to take longer as LRU is the same FIFO queue, but with an extra table for book keeping.

What surprised me was the two different random number generation algorithms taking nearly the same amount of time. I expected for the mersenne twister to take longer than the LCG due to the complexity differences between them. However, the LCG took around 0s to 0.1s longer in every case.