

CS 433 (2) HW5 Report

Boris Nikulin

2019-05-11

Contents

1 Data Analysis	1
1.1 Data Import	1
1.2 Data Analysis	2
1.3 Data Visualization	6

1 Data Analysis

1.1 Data Import

First we load the simulation data, filter out processes that never finished, and give the data a once over.

However, using `readr::read_tsv` to load the data take's minutes because the data per run is around 1 GB. After all the runs of the simulation, the final dataset is 4.2 GB. As a side note, the gzip comprised size of the whole data set is 444 MB. For this reason, we will primarily use `data.table` instead of `readr` and `dplyr`. `data.table::fread` reads the whole dataset in 10 s to 20 s.

```
library(data.table)
library(dplyr)

##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:data.table':
##
##   between, first, last
## The following objects are masked from 'package:stats':
##
##   filter, lag
## The following objects are masked from 'package:base':
##
```

```
## intersect, setdiff, setequal, union
library(magrittr)

data <- fread('paging_sim.tsv')

data %<>% .[, time := time / 1E6]

glimpse(data)
## Observations: 119,364,413
## Variables: 7
## $ time <dbl> 0.0e+00, 1.0e-06, 1.6e-05, 1.8...
## $ num_pages <int> 524288, 524288, 524288, 524288...
## $ page_size_bytes <int> 256, 256, 256, 256, 256, ...
## $ memory_size_bytes <int> 65536, 65536, 65536, 65536, 65...
## $ replacement_algorithm <chr> "fifo", "fifo", "fifo", "fifo"...
## $ event <chr> "reference", "fault", "referen...
## $ page <int> 397863, 397863, 397863, 281415...
```

The times are in seconds.

1.2 Data Analysis

```
run_group <- c(
  'num_pages',
  'page_size_bytes',
  'memory_size_bytes',
  'replacement_algorithm'
)

(data_time <- data[, .(time_max = max(time)), mget(run_group)])

##      num_pages page_size_bytes memory_size_bytes
## 1:      524288           256           65536
## 2:      524288           256           65536
## 3:      524288           256           65536
## 4:      524288           256           65536
## 5:      524288           256          16777216
## 6:      524288           256          16777216
## 7:      524288           256          16777216
## 8:      524288           256          16777216
## 9:       16384          8192           65536
## 10:      16384          8192           65536
## 11:      16384          8192           65536
## 12:      16384          8192           65536
## 13:      16384          8192          16777216
```

```

## 14:      16384      8192      16777216
## 15:      16384      8192      16777216
## 16:      16384      8192      16777216
##      replacement_algorithm  time_max
## 1:              fifo 14.284713
## 2:      rng:mt19937  6.528269
## 3:      rng:minstd_rand  6.686142
## 4:              lru 16.293392
## 5:              fifo  7.209118
## 6:      rng:mt19937  6.364927
## 7:      rng:minstd_rand  6.473702
## 8:              lru  9.296487
## 9:              fifo 14.972721
## 10:      rng:mt19937  5.647755
## 11:      rng:minstd_rand  5.795382
## 12:              lru 16.359806
## 13:              fifo  6.653178
## 14:      rng:mt19937  5.742128
## 15:      rng:minstd_rand  5.787111
## 16:              lru  8.094985

(data_event_count <- data[, .N, mget(append(run_group, 'event'))])

##      num_pages page_size_bytes memory_size_bytes
## 1:      524288      256      65536
## 2:      524288      256      65536
## 3:      524288      256      65536
## 4:      524288      256      65536
## 5:      524288      256      65536
## 6:      524288      256      65536
## 7:      524288      256      65536
## 8:      524288      256      65536
## 9:      524288      256      65536
## 10:      524288      256      65536
## 11:      524288      256      65536
## 12:      524288      256      65536
## 13:      524288      256      65536
## 14:      524288      256      65536
## 15:      524288      256      65536
## 16:      524288      256      65536
## 17:      524288      256      16777216
## 18:      524288      256      16777216
## 19:      524288      256      16777216
## 20:      524288      256      16777216
## 21:      524288      256      16777216
## 22:      524288      256      16777216

```

## 23:	524288	256	16777216
## 24:	524288	256	16777216
## 25:	524288	256	16777216
## 26:	524288	256	16777216
## 27:	524288	256	16777216
## 28:	524288	256	16777216
## 29:	524288	256	16777216
## 30:	524288	256	16777216
## 31:	524288	256	16777216
## 32:	524288	256	16777216
## 33:	16384	8192	65536
## 34:	16384	8192	65536
## 35:	16384	8192	65536
## 36:	16384	8192	65536
## 37:	16384	8192	65536
## 38:	16384	8192	65536
## 39:	16384	8192	65536
## 40:	16384	8192	65536
## 41:	16384	8192	65536
## 42:	16384	8192	65536
## 43:	16384	8192	65536
## 44:	16384	8192	65536
## 45:	16384	8192	65536
## 46:	16384	8192	65536
## 47:	16384	8192	65536
## 48:	16384	8192	65536
## 49:	16384	8192	16777216
## 50:	16384	8192	16777216
## 51:	16384	8192	16777216
## 52:	16384	8192	16777216
## 53:	16384	8192	16777216
## 54:	16384	8192	16777216
## 55:	16384	8192	16777216
## 56:	16384	8192	16777216
## 57:	16384	8192	16777216
## 58:	16384	8192	16777216
## 59:	16384	8192	16777216
## 60:	16384	8192	16777216
## 61:	16384	8192	16777216
## 62:	16384	8192	16777216
## 63:	16384	8192	16777216
## 64:	16384	8192	16777216
##	num_pages	page_size_bytes	memory_size_bytes
##	replacement_algorithm	event	N
## 1:	fifo	reference	5000000

```

## 2:          fifo      fault 2778912
## 3:          fifo      replace 2778656
## 4:          fifo      flush 2483210
## 5:          rng:mt19937 reference 5000000
## 6:          rng:mt19937      fault 336407
## 7:          rng:mt19937      replace 336151
## 8:          rng:mt19937      flush 61818
## 9:          rng:minstd_rand reference 5000000
## 10:         rng:minstd_rand      fault 336087
## 11:         rng:minstd_rand      replace 335831
## 12:         rng:minstd_rand      flush 61562
## 13:         lru reference 5000000
## 14:         lru      fault 2632039
## 15:         lru      replace 2631783
## 16:         lru      flush 2336294
## 17:         fifo reference 5000000
## 18:         fifo      fault 514905
## 19:         fifo      replace 449369
## 20:         fifo      flush 300719
## 21:         rng:mt19937 reference 5000000
## 22:         rng:mt19937      fault 332699
## 23:         rng:mt19937      replace 267163
## 24:         rng:mt19937      flush 58969
## 25:         rng:minstd_rand reference 5000000
## 26:         rng:minstd_rand      fault 332344
## 27:         rng:minstd_rand      replace 266808
## 28:         rng:minstd_rand      flush 58657
## 29:         lru reference 5000000
## 30:         lru      fault 482163
## 31:         lru      replace 416627
## 32:         lru      flush 266250
## 33:         fifo reference 5000000
## 34:         fifo      fault 2965735
## 35:         fifo      replace 2965727
## 36:         fifo      flush 2943957
## 37:         rng:mt19937 reference 5000000
## 38:         rng:mt19937      fault 31530
## 39:         rng:mt19937      replace 31522
## 40:         rng:mt19937      flush 16567
## 41:         rng:minstd_rand reference 5000000
## 42:         rng:minstd_rand      fault 31487
## 43:         rng:minstd_rand      replace 31479
## 44:         rng:minstd_rand      flush 16468
## 45:         lru reference 5000000
## 46:         lru      fault 2663297

```


```
## 47:          lru   replace 2663289
## 48:          lru    flush 2641522
## 49:         fifo reference 5000000
## 50:         fifo   fault  249038
## 51:         fifo   replace 246990
## 52:         fifo    flush  233736
## 53:      rng:mt19937 reference 5000000
## 54:      rng:mt19937   fault   31497
## 55:      rng:mt19937   replace  29449
## 56:      rng:mt19937    flush  16505
## 57:  rng:minstd_rand reference 5000000
## 58:  rng:minstd_rand   fault   31249
## 59:  rng:minstd_rand   replace  29201
## 60:  rng:minstd_rand    flush  16314
## 61:          lru reference 5000000
## 62:          lru   fault  213375
## 63:          lru   replace 211327
## 64:          lru    flush  197729
##      replacement_algorithm   event      N
data_event_count %<>% .[event != 'reference']
```

Another side note, `awk` takes around 10s per simulation to generate the counts while `data.table` gets all the counts at once in under 10s.

1.3 Data Visualization

```
library(ggplot2)

ggplot(data_event_count,
       aes(event, N, fill = replacement_algorithm)) +
  geom_col(position = 'dodge') +
  facet_grid(memory_size_bytes ~ page_size_bytes) +
  labs(x = 'Event', y = 'Count', fill = 'Algorithm')
```



rng:mt19937 refers to the C++ `random` library and is a 32 bit mersenne twister default constructed. rng:minstd_rand also corresponds to the `random` library and is the newer 1993 minimum standard LCG. It is also default constructed.


The bar graph above has some key takeaways. Page size, shown in the horizontal facets, does not reduce the event counts nearly as much as memory size, shown in the vertical facets. At the end of the day, more memory is best.

Another takeaway is that when there is insufficient memory, page size has no effect or a small adverse affect on the number of events excluding the random algorithm. Page size does have an effect when memory size is not the limiting factor.

Lastly, random page replacement is unfairly good for how simple the algorithm is and is primarily affected by page size unlike the other algorithms.

```
ggplot(data_time, aes(replacement_algorithm, time_max)) +
  geom_col() +
```

```
facet_grid(memory_size_bytes ~ page_size_bytes) +
labs(x = 'Algorithm', y = 'Time (s)') +
coord_flip()
```



The times for the final events, which included the time to log each event, are pretty similar.

I expected LRU to take longer as LRU is the same FIFO queue, but with an extra table for book keeping.

What surprised me was the two different random number generation algorithms taking nearly the same amount of time. I expected for the mersenne twister to take longer than the LCG due to the complexity differences between them. However, the LCG took around 0s to 0.1s longer in every case.